


Name: _____ Timur Guner _____

Wireshark Lab: IP v7.0

Supplement to *Computer Networking: A Top-Down Approach*, 7th ed., J.F. Kurose and K.W. Ross

“Tell me and I forget. Show me and I remember. Involve me and I understand.” Chinese proverb

© 2005-2016, J.F Kurose and K.W. Ross, All Rights Reserved



In this lab, we'll investigate the IP protocol, focusing on the IP datagram. We'll do so by analyzing a trace of IP datagrams sent and received by an execution of the `traceroute` program (the `traceroute` program itself is explored in more detail in the Wireshark ICMP lab). We'll investigate the various fields in the IP datagram, and study IP fragmentation in detail.

Before beginning this lab, you'll probably want to review sections 1.4.3 in the text¹ and section 3.4 of RFC 2151 [[ftp://ftp.rfc-editor.org/in-notes/rfc2151.txt](http://ftp.rfc-editor.org/in-notes/rfc2151.txt)] to update yourself on the operation of the `traceroute` program. You'll also want to read Section 4.3 in the text, and probably also have RFC 791 [[ftp://ftp.rfc-editor.org/in-notes/rfc791.txt](http://ftp.rfc-editor.org/in-notes/rfc791.txt)] on hand as well, for a discussion of the IP protocol.

1. Capturing packets from an execution of `traceroute`

In order to generate a trace of IP datagrams for this lab, we'll use the `traceroute` program to send datagrams of different sizes towards some destination, *X*. Recall that `traceroute` operates by first sending one or more datagrams with the time-to-live (TTL) field in the IP header set to 1; it then sends a series of one or more datagrams towards the same destination with a TTL value of 2; it then sends a series of datagrams towards the same destination with a TTL value of 3; and so on. Recall that a router must decrement the TTL in each received datagram by 1 (actually, RFC 791 says that the router must decrement the TTL by *at least* one). If the TTL reaches 0, the router returns an ICMP message (type 11 – TTL-exceeded) to the sending host. As a result of this behavior, a datagram with a TTL of 1 (sent by the host executing `traceroute`) will cause the router one hop away from the sender to send an ICMP TTL-exceeded message back to the sender; the datagram sent with a TTL of 2 will cause the router two hops

¹ References to figures and sections are for the 7th edition of our text, *Computer Networks, A Top-down Approach*, 7th ed., J.F. Kurose and K.W. Ross, Addison-Wesley/Pearson, 2016.

away to send an ICMP message back to the sender; the datagram sent with a TTL of 3 will cause the router three hops away to send an ICMP message back to the sender; and so on. In this manner, the host executing `tracert` can learn the identities of the routers between itself and destination X by looking at the source IP addresses in the datagrams containing the ICMP TTL-exceeded messages.

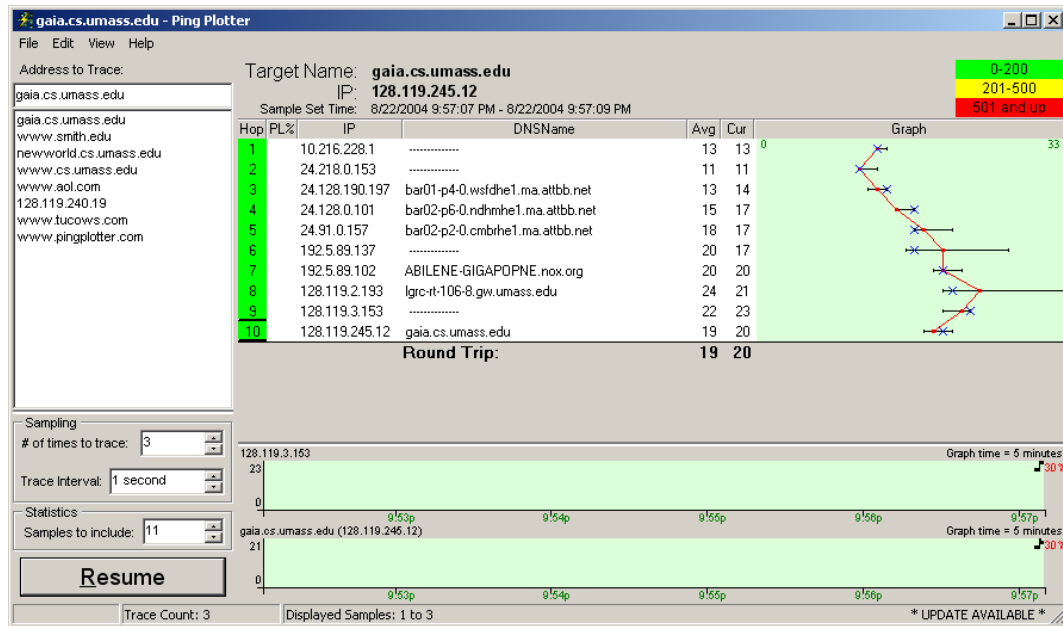
We'll want to run `tracert` and have it send datagrams of various lengths.

- **Windows.** The `tracert` program (used for our ICMP Wireshark lab) provided with Windows does not allow one to change the size of the ICMP echo request (ping) message sent by the `tracert` program. A nicer Windows `tracert` program is *pingplotter*, available both in free version and shareware versions at <http://www.pingplotter.com>. Download and install *pingplotter*, and test it out by performing a few traceroutes to your favorite sites. The size of the ICMP echo request message can be explicitly set in *pingplotter* by selecting the menu item *Edit->Options->Packet Options* and then filling in the *Packet Size* field. The default packet size is 56 bytes. Once *pingplotter* has sent a series of packets with the increasing TTL values, it restarts the sending process again with a TTL of 1, after waiting *Trace Interval* amount of time. The value of *Trace Interval* and the number of intervals can be explicitly set in *pingplotter*.
- **Linux/Unix/MacOS.** With the Unix/MacOS `tracert` command, the size of the UDP datagram sent towards the destination can be explicitly set by indicating the number of bytes in the datagram; this value is entered in the `tracert` command line immediately after the name or address of the destination. For example, to send `tracert` datagrams of 2000 bytes towards `gaia.cs.umass.edu`, the command would be:

```
%tracert gaia.cs.umass.edu 2000
```

Do the following:

- Start up Wireshark and begin packet capture (*Capture->Start*) and then press *OK* on the Wireshark Packet Capture Options screen (we'll not need to select any options here).
- If you are using a Windows platform, start up *pingplotter* and enter the name of a target destination in the "Address to Trace Window." Enter 3 in the "# of times to Trace" field, so you don't gather too much data. Select the menu item *Edit->Advanced Options->Packet Options* and enter a value of 56 in the *Packet Size* field and then press *OK*. Then press the *Trace* button. You should see a *pingplotter* window that looks something like this:



Next, send a set of datagrams with a longer length, by selecting *Edit->Advanced Options->Packet Options* and enter a value of 2000 in the *Packet Size* field and then press OK. Then press the Resume button.

Finally, send a set of datagrams with a longer length, by selecting *Edit->Advanced Options->Packet Options* and enter a value of 3500 in the *Packet Size* field and then press OK. Then press the Resume button.

Stop Wireshark tracing.

- If you are using a Unix or Mac platform, enter three `traceroute` commands, one with a length of 56 bytes, one with a length of 2000 bytes, and one with a length of 3500 bytes.

Stop Wireshark tracing.

If you are unable to run Wireshark on a live network connection, you can download a packet trace file that was captured while following the steps above on one of the author's Windows computers². You may well find it valuable to download this trace even if you've captured your own trace and use it, as well as your own trace, when you explore the questions below.

² Download the zip file <http://gaia.cs.umass.edu/wireshark-labs/wireshark-traces.zip> and extract the file *ip-ethereal-trace-1*. The traces in this zip file were collected by Wireshark running on one of the author's computers, while performing the steps indicated in the Wireshark lab. Once you have downloaded the trace, you can load it into Wireshark and view the trace using the *File* pull down menu, choosing *Open*, and then selecting the *ip-ethereal-trace-1* trace file.

1. Select the first ICMP Echo Request message sent by your computer, and expand the Internet Protocol part of the packet in the packet details window.
What is the IP address of your computer?
Circled in red is the IP address is 10.0.0.172
2. Within the IP packet header, what is the value in the upper layer protocol field?
Circled in blue is the value is 0x0001

3. How many bytes are in the IP header? How many bytes are in the payload *of the IP datagram*? Explain how you determined the number of payload bytes.
The header is 20 bytes. Total length is 56 so the payload is 56-20=36 bytes. Check the values circled in black
4. Has this IP datagram been fragmented? Explain how you determined whether or not the datagram has been fragmented.
The IP has not been fragmented because fragment offset is 0 and the more fragments bit is 0. Highlight shows the more fragment offset and right above it the more fragments set to 0 as well.

Next, sort the traced packets according to IP source address by clicking on the *Source* column header; a small downward pointing arrow should appear next to the word *Source*. If the arrow points up, click on the *Source* column header again. Select the first ICMP Echo Request message sent by your computer, and expand the Internet Protocol portion in the “details of selected packet header” window. In the “listing of captured packets” window, you should see all of the subsequent ICMP messages (perhaps with additional interspersed packets sent by other protocols running on your computer) below this first ICMP. Use the down arrow to move through the ICMP messages sent by your computer.

5. Which fields in the IP datagram *always* change from one datagram to the next within this series of ICMP messages sent by your computer?
Identification, checksum, and time to live
6. Which fields stay constant? Which of the fields *must* stay constant? Which fields must change? Why?

The fields that stay constant across the IP datagrams are:

- Version
- Header length
- Source IP
- Destination IP
- Differentiated Services
- Upper Layer Protocol (since these are ICMP packets)

The fields that must stay constant are:

- Same as above

The fields that must change are:

- Identification because the packet ids need to be different
- Time to live increments with each packet
- Header checksum since the header changes, the checksum also changes

7. Describe the pattern you see in the values in the Identification field of the IP datagram

The identification field is increasing by one each time. First request started at 0xee91.

Next (with the packets still sorted by source address) find the series of ICMP TTL-exceeded replies sent to your computer by the nearest (first hop) router.

8. What is the value in the Identification field and the TTL field?

114	9.008331	10.0.0.172	128.119.245.12	ICMP	70 Echo (ping) request	id=0x0001, seq=9095/34595, ttl=1 (no response found!)
115	9.014426	10.0.0.1	10.0.0.172	ICMP	98 Time-to-live exceeded	(Time to live exceeded in transit)
118	9.047357	10.0.0.172	128.119.245.12	ICMP	70 Echo (ping) request	id=0x0001, seq=9096/34851, ttl=2 (no response found!)
120	9.064815	128.119.245.12	10.0.0.172	ICMP	70 Echo (ping) reply	id=0x0001, seq=9094/34339, ttl=33 (request in 112)
121	9.069763	10.60.140.2	10.0.0.172	ICMP	70 Time-to-live exceeded	(Time to live exceeded in transit)
126	9.085054	10.0.0.1	10.0.0.172	ICMP	120 Destination unreachable	(Port unreachable)
127	9.085696	10.0.0.172	128.119.245.12	ICMP	70 Echo (ping) request	id=0x0001, seq=9097/35107, ttl=3 (no response found!)
129	9.098139	96.108.161.45	10.0.0.172	ICMP	98 Time-to-live exceeded	(Time to live exceeded in transit)
132	9.125206	10.0.0.172	128.119.245.12	ICMP	70 Echo (ping) request	id=0x0001, seq=9098/35363, ttl=4 (no response found!)
134	9.136741	68.85.155.161	10.0.0.172	ICMP	98 Time-to-live exceeded	(Time to live exceeded in transit)
136	9.163326	10.0.0.172	128.119.245.12	ICMP	70 Echo (ping) request	id=0x0001, seq=9099/35619, ttl=5 (no response found!)
138	9.176350	68.87.192.249	10.0.0.172	ICMP	110 Time-to-live exceeded	(Time to live exceeded in transit)
140	9.203065	10.0.0.172	128.119.245.12	ICMP	70 Echo (ping) request	id=0x0001, seq=9100/35875, ttl=6 (no response found!)
142	9.219642	162.151.86.57	10.0.0.172	ICMP	110 Time-to-live exceeded	(Time to live exceeded in transit)
144	9.241185	10.0.0.172	128.119.245.12	ICMP	70 Echo (ping) request	id=0x0001, seq=9101/36131, ttl=7 (no response found!)


```

Frame 115: 98 bytes on wire (784 bits), 98 bytes captured (784 bits) on interface \Device\NPF_{F1521ED5-A867-4238-A0A7-589345A46658}, id 0
Ethernet II, Src: dl:b5:d8:cc:56:a4 (dl:b5:d8:cc:56:a4), Dst: IntelCor_66:59:f4 (08:71:90:66:59:f4)
Internet Protocol Version 4, Src: 10.0.0.1, Dst: 10.0.0.172
  0100 .... = Version: 4
  .... 0101 = Header Length: 20 bytes (5)
  Differentiated Services Field: 0xc0 (DSCP: CS6, ECN: Not-ECT)
    1100 00.. = Differentiated Services Codepoint: Class Selector 6 (48)
    .... 00.. = Explicit Congestion Notification: Not ECN-Capable Transport (0)
  Total Length: 84
  Identification: 0x248e (9358)
  Flags: 0x00
    0... .... = Reserved bit: Not set
    .0.. .... = Don't fragment: Not set
    ..0. .... = More fragments: Not set
  ...0 0000 0000 0000 = Fragment Offset: 0
  Time to Live: 64
  Protocol: ICMP (1)
  Header Checksum: 0x40af [validation disabled]
  [Header checksum status: Unverified]
  Source Address: 10.0.0.1
  Destination Address: 10.0.0.172
Internet Control Message Protocol
  Type: 11 (Time-to-live exceeded)
  Code: 0 (Time to live exceeded in transit)

```

Identification number is 0x248e and the TTL is 64

9. Do these values remain unchanged for all of the ICMP TTL-exceeded replies sent to your computer by the nearest (first hop) router? Why?

The identification field will change for all replies because it is a unique value. When two or more datagrams have the same id, it means it is fragmented.

The TTL does not change because it is the same for all first hops

Fragmentation

Sort the packet listing according to time again by clicking on the *Time* column.

10. Find the first ICMP Echo Request message that was sent by your computer after you changed the *Packet Size* in *pingplotter* to be 2000. Has that message been fragmented across more than one IP datagram?

26	2.765150	192.168.1.133	255.255.255.255	UDP	65 3073 → 6524 Len=23
27	2.765150	192.168.1.133	255.255.255.255	UDP	65 3073 → 6524 Len=23
28	2.765150	192.168.1.133	255.255.255.255	UDP	65 3074 → 35344 Len=23
29	2.765150	192.168.1.133	255.255.255.255	UDP	65 3074 → 35344 Len=23
30	3.070322	SnapAV_23:1f:ad	Broadcast	ARP	60 Who has 192.168.1.122? Tell 192.168.1.19
31	3.379440	192.168.1.132	255.255.255.255	UDP	65 37553 → 6524 Len=23
32	3.379440	192.168.1.132	255.255.255.255	UDP	65 35408 → 35344 Len=23
33	3.379440	192.168.1.132	255.255.255.255	UDP	65 37553 → 6524 Len=23
34	3.379440	192.168.1.132	255.255.255.255	UDP	65 35408 → 35344 Len=23
35	3.381854	192.168.1.86	192.168.1.255	UDP	1066 10102 → 10102 Len=1024
36	3.381854	192.168.1.17	239.255.255.250	UDP	308 34235 → 1902 Len=266
37	3.381971	192.168.1.103	192.168.1.255	UDP	1066 10102 → 10102 Len=1024
38	3.468790	Ubiquiti_cd:83:4d	IntelCor_66:59:f4	ARP	60 Who has 192.168.1.7? Tell 192.168.1.1
39	3.468826	IntelCor_66:59:f4	Ubiquiti_cd:83:4d	ARP	42 192.168.1.7 is at 08:71:90:66:59:f4
40	3.542771	192.168.1.7	128.119.245.12	IPv4	1514 Fragmented IP protocol (proto=ICMP 1, off=0, ID=7942) [Reassembled in #42]
41	3.542771	192.168.1.7	128.119.245.12	IPv4	1514 Fragmented IP protocol (proto=ICMP 1, off=1480, ID=7942) [Reassembled in #42]
42	3.542771	192.168.1.7	128.119.245.12	ICMP	554 Echo (ping) request id=0x0001, seq=11832/14382, ttl=255 (reply in #4)
43	3.581028	192.168.1.7	128.119.245.12	IPv4	1514 Fragmented IP protocol (proto=ICMP 1, off=0, ID=7943) [Reassembled in #45]
44	3.581028	192.168.1.7	128.119.245.12	IPv4	1514 Fragmented IP protocol (proto=ICMP 1, off=1480, ID=7943) [Reassembled in #45]
45	3.581028	192.168.1.7	128.119.245.12	ICMP	554 Echo (ping) request id=0x0001, seq=11833/14638, ttl=1 (no response found!)
46	3.584133	192.168.1.1	192.168.1.7	ICMP	590 Time-to-live exceeded (Time to live exceeded in transit)
47	3.620909	192.168.1.7	128.119.245.12	IPv4	1514 Fragmented IP protocol (proto=ICMP 1, off=0, ID=7944) [Reassembled in #49]
48	3.620909	192.168.1.7	128.119.245.12	IPv4	1514 Fragmented IP protocol (proto=ICMP 1, off=1480, ID=7944) [Reassembled in #49]
49	3.620909	192.168.1.7	128.119.245.12	ICMP	554 Echo (ping) request id=0x0001, seq=11834/14894, ttl=2 (no response found!)
50	3.622134	192.168.1.7	192.168.1.1	DNS	84 Standard query 0xe477 PTR 1.1.168.192.in-addr.arpa
51	3.623896	192.168.1.1	192.168.1.7	DNS	105 Standard query response 0xe477 PTR 1.1.168.192.in-addr.arpa PTR Gateway
52	3.629870	128.119.245.12	192.168.1.7	IPv4	1514 Fragmented IP protocol (proto=ICMP 1, off=0, ID=6562) [Reassembled in #54]
53	3.629870	128.119.245.12	192.168.1.7	IPv4	1514 Fragmented IP protocol (proto=ICMP 1, off=1480, ID=6562) [Reassembled in #54]
54	3.629870	128.119.245.12	192.168.1.7	ICMP	554 Echo (ping) reply id=0x0001, seq=11832/14382, ttl=46 (request in #42)
55	3.652881	10.100.12.1	192.168.1.7	ICMP	590 Time-to-live exceeded (Time to live exceeded in transit)
56	3.659852	192.168.1.7	128.119.245.12	IPv4	1514 Fragmented IP protocol (proto=ICMP 1, off=0, ID=7945) [Reassembled in #58]

```

<
> Frame 42: 554 bytes on wire (4432 bits), 554 bytes captured (4432 bits) on interface \Device\NPF_{F1521ED5-A867-4238-ABA7-589345A46658}, id 0
> Ethernet II, Src: IntelCor_66:59:f4 (08:71:90:66:59:f4), Dst: Ubiquiti_cd:83:4d (08:2a:a8:cd:83:4d)
✓ Internet Protocol Version 4, Src: 192.168.1.7, Dst: 128.119.245.12
  0100 .... = Version: 4
  .... 0101 = Header Length: 20 bytes (5)
  > Differentiated Services Field: 0x00 (DSCP: CS0, ECN: Not-ECT)
  Total Length: 540
  Identification: 0x7942 (31042)
  ✓ Flags: 0x01
    0... .... = Reserved bit: Not set
    .0... .... = Don't fragment: Not set
    ..0... .... = More fragments: Not set
    ...0 1011 1001 0000 = Fragment Offset: 2960
  Time to Live: 255
  Protocol: ICMP (1)
  Header Checksum: 0x0000 [validation disabled]
  [Header checksum status: Unverified]
  Source Address: 192.168.1.7
  Destination Address: 128.119.245.12
  > [3 IPv4 Fragments (3480 bytes): #40(1480), #41(1480), #42(520)]
✓ Internet Control Message Protocol
  Type: 8 (Echo (ping) request)
  Code: 0
  Checksum: 0xefec [correct]
  [Checksum Status: Good]
  Identifier (BE): 1 (0x0001)
  Identifier (LE): 256 (0x0100)
  Sequence Number (BE): 11832 (0x2e38)
  Sequence Number (LE): 11832 (0x2e38)

```

14. How many fragments were created from the original datagram?

Three fragments were created from the original data, as shown in the above screenshot.

15. What fields change in the IP header among the fragments?

The fields that change through all three is the fragment offset (0, 1480, 2960) and the header checksum. The first two fragments have a length of 1500 bytes and the more fragments bit is set. The last fragment has a length of 540 bytes and the more fragment bit is set to 0.